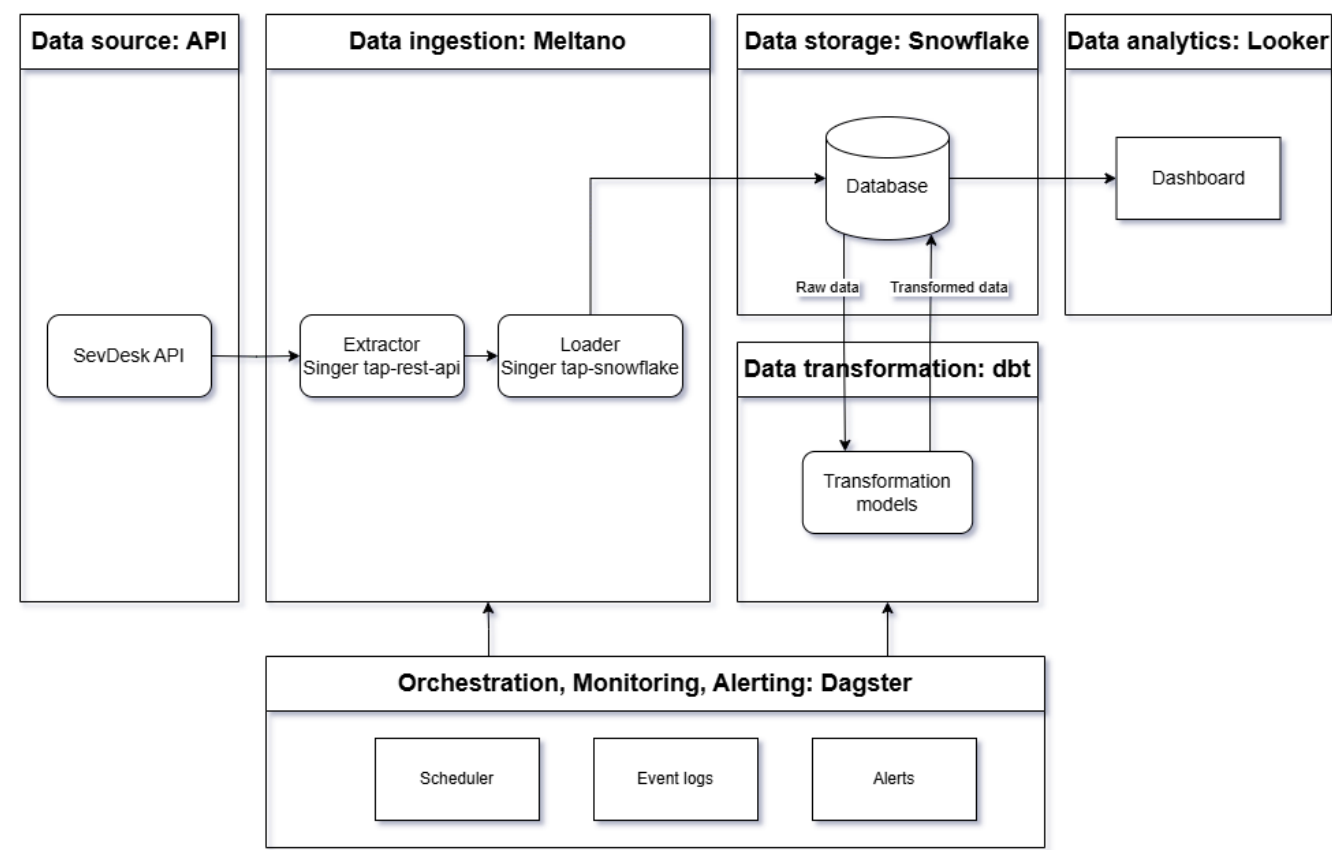


Exercise: Build a Data Platform Design

System Architecture Diagram

SevDesk System Architecture Diagram



Tech Stack

Component	Tool	Purpose	Why?
Orchestration	Dagster	Schedule, execute, and monitor pipelines (Meltano + dbt workflows)	Modern orchestrator with lineage, status, retriy policies
Ingestion	Meltano SDK	Uses Singer taps (APlextractor + Snowflake loader)	Plugin-based extract/load, flexible to create custom plugins
Data Warehouse	Snowflake	Central storage of raw + transformed data	Cloud-native DWH with scalability, security, and dbt compatibility
Transformation	dbt	SQL-based modeling and transformations	Integrates with Snowflake, git-friendly, future-ready for Looker
Analytics	Looker	BI dashboard and exploration layer	Connects directly to Snowflake, supports dbt layer

Component	Tool	Purpose	Why?
Monitoring & Alerting	Dagster	Observability and alerting	UI for visual monitoring of logs, built-in alerting system (Slack/Email)

Data Pipeline Steps for SevDesk invoices

1. Data extraction

- Job: set and scheduled on Dagster
- Frequency: the meltano pipeline is run on a fixed schedule (e.g., every hour)
- Extractor: `tap-rest-api-msdk` in Meltano.
- Configuration:
 - API URL: `https://my.sevdesk.de/api/v1/`
 - Headers: `Authorization header ({ env('SEVDESK_API_KEY') })`, dynamically set from environment variables stored in `.env`
 - Streams: The pipeline is designed to extract objects from `/Invoice` endpoint

```
meltano.yml
...
extractors:
- name: tap-rest-api-msdk
  variant: widen
  pip_url: tap-rest-api-msdk
  config:
    api_url: "https://my.sevdesk.de/api/v1"
    headers:
      Authorization: "{ env('SEVDESK_API_KEY') }"

  streams:
  - name: "invoice"
    path: "/Invoice"
    method: "GET"
    records_path: "objects"
...
```

2. Data loading

- Loader: `tap-snowflake` in Meltano
- Configuration:

```
meltano.yml
...
loaders:
- name: target-snowflake
  variant: meltanolabs
  pip_url: meltanolabs-target-snowflake
```

```
config:
  user: "{{ env_var('SNOWFLAKE_USER') }}"
  password: "{{ env_var('SNOWFLAKE_PASSWORD') }}"
  account: "{{ env_var('SNOWFLAKE_ACCOUNT') }}"
  warehouse: "{{ env_var('SNOWFLAKE_WAREHOUSE') }}"
  database: "{{ env_var('SNOWFLAKE_DATABASE') }}"
  schema: "{{ env_var('SNOWFLAKE_SCHEMA') }}"
  role: "{{ env_var('SNOWFLAKE_ROLE') }}"
...
```

- Output: load data into **SEVDESK.raw.invoices** table in Snowflake

3. Data Transformation

- Trigger: Dagster triggers the dbt transformation jobs to run directly on Snowflake
- Transformation: dbt models will clean and reshape this raw data to meet business requirements. The transformations can include aggregations, filtering, or creating new dimensions
- Output: cleaned data into **SEVDESK.staging.invoices_stg**, curated data into **analytics.invoices_summary** for e.g (business-ready models)

4. Data analytics

- Looker Integration: Looker can directly query **analytics.invoices_summary** in Snowflake for business intelligence and analytics reporting
- Output: insight extraction from the data into dashboard

5. Monitoring and alerting

- Visualization: job executions can be accessed real time on Dagster UI
- Configuration: send Slack/email notification to stakeholders whenever a pipeline fails, or retries are needed

Use Case: Invoice Revenue Ranking by Country

The Finance or Sales team wants a dashboard that shows:

- The total invoice amount per city in €
- Updated every hour

1. Invoice data is ingested using maltano extractor triggered every hour by dagster

2. Data is loaded into snowflake raw schema

SEVDESK / RAW / INVOICES

Table ACCOUNTADMIN 1 hour ago 0 0.0B

Table DetailsColumnsData PreviewCopy History

COMPUTE_WH2 Rows • Updated just now

ID	OBJECTNAME	INVOICENUMBER	CONTACT_ID	create	update
194036284	Invoice	RE-1000	null	2025-04-06 18:33:03.000	2025-04-06 18:33:03.000

3. dbt staging models clean fields, cast types, convert ammount to the same currency, create a new table `analytics.invoices_amount_by_city`

Home

Search

Projects

Data

Databases

Add Data

Migrations

Data Products

AI & ML

Monitoring

Admin

SEVDESK

ANALYTICS

Tables

INVOICES_AMOUNT_BY_CITY

INVOICES_SUMMARY

DBT

Tables

ENRICHED

INFORMATION_SCHEMA

Views

PUBLIC

No Objects found

RAW

Tables

INVOICES

STAGING

Tables

INVOICES_STG

4. Looker query the table in snowflake to generate a dashboard

5. Dagster sends a notification in case one of the pipeline failed so that the finance team knows the dashboard is outdated

Possible Enhancements

Incremental Loading

- Problem:
Loading all data every time from the SevDesk API can be inefficient, especially as the data grows over time
- Solution:
Implement incremental loading using **update** timestamps in the SevDesk API. Only new or updated records are loaded, reducing the amount of data transferred and stored.
- Implementation:
Modify the extractor (tap-rest-api-msdk) to fetch data based on **update** fields

Event-Driven Ingestion via Webhooks

- Problem:
Relying on a fixed schedule for data extraction can result in delays
- Solution: Integrate webhooks to trigger data extraction whenever there is new data in SevDesk (e.g., a new invoice or payment)
- Implementation: Implement a webhook listener to trigger the Meltano extraction pipeline upon receipt of an event from SevDesk. This can be done by using an HTTP API in Dagster

CI/CD for Pipeline + dbt Deployment

- Problem:
Maintaining pipeline and transformation code manually can be error-prone
- Solution:
Use a CI/CD pipeline to automatically deploy and test changes to both the Meltano pipeline and dbt transformations
- Implementation:
Set up a CI/CD pipeline with GitHub Actions to run tests and deploy updates automatically when code is pushed to the repository